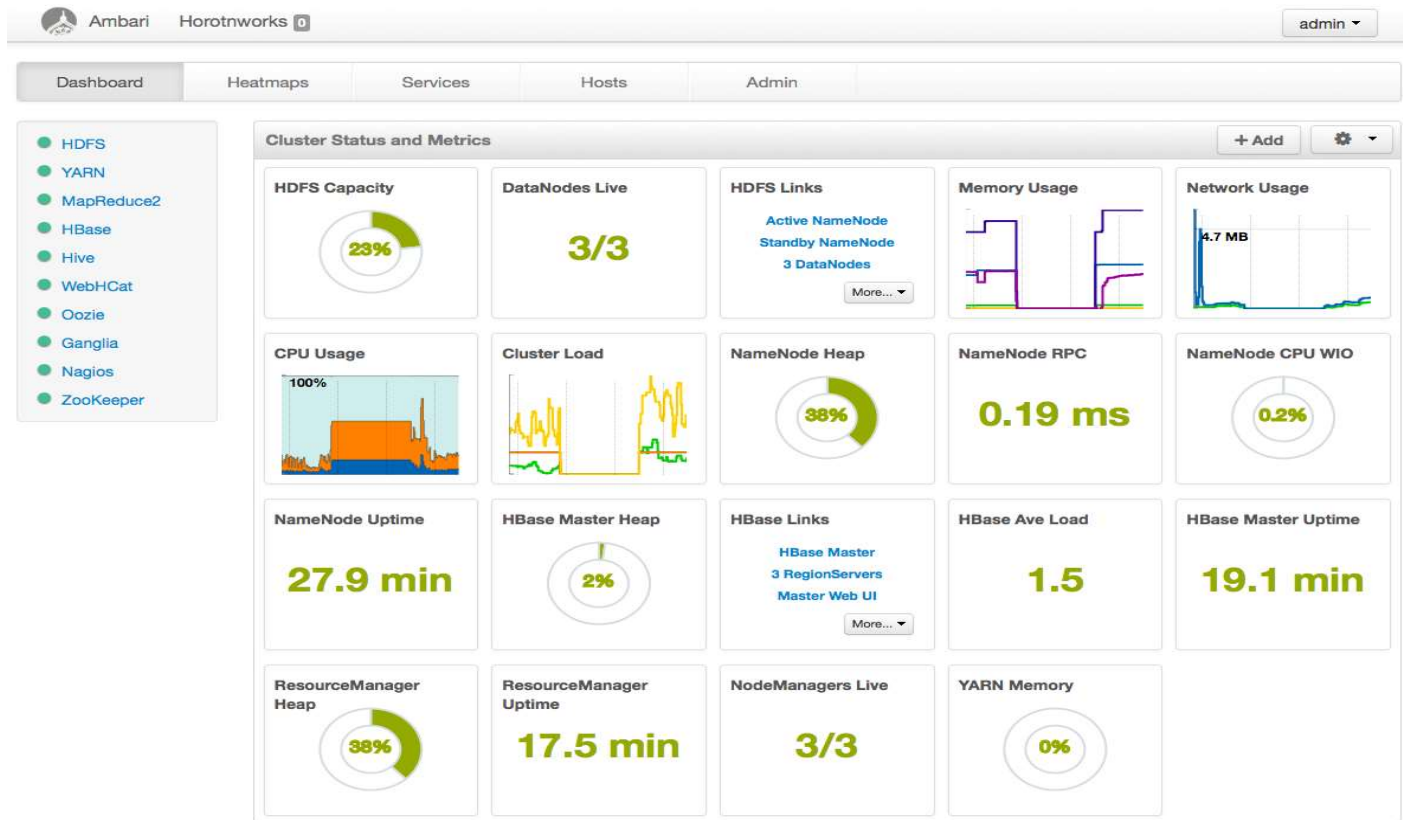


# Optimizing MapReduce Jobs

# Monitoring Performance






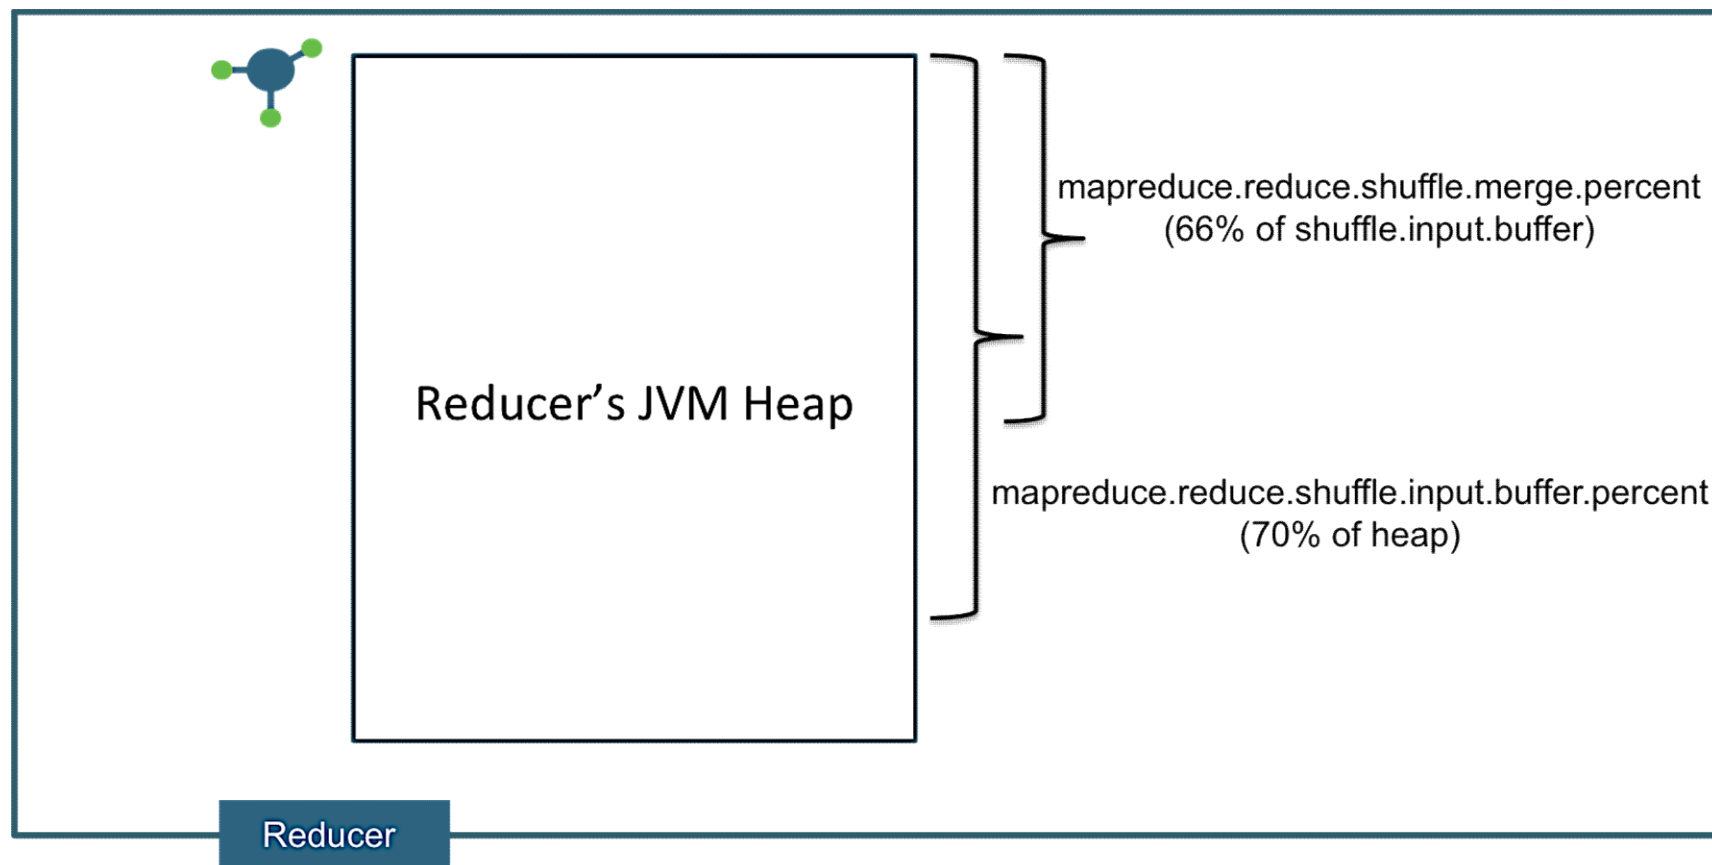
# Optimization Best Practices

- Distribute workload evenly across NodeManagers
- Use a combiner
- Avoid instantiating new objects, including string literals
- Prefer StringBuilder over string concatenation
- Use data compression
- Avoid converting numeric types to text
- Define and configure a RawComparator
- Prefer StringUtils.split over String.split

# Optimizing the Map Phase



# Optimizing the Reduce Phase



# Data Compression

- **Snappy**: `org.apache.hadoop.io.compress.SnappyCodec`
- **gzip**: `org.apache.hadoop.io.compress.GzipCodec`
- **bzip2**: `org.apache.hadoop.io.compress.BZip2Codec`
- **LZO**: `com.hadoop.compression.lzo.LzopCodec`
- **DEFLATE**: `org.apache.hadoop.io.compress.DefaultCodec`

## Limitations of Compression

- **Space vs. time:** is the gain in filesystem space or smaller network traffic worth the additional time it takes to compress and decompress the data?
- **Splittable vs. Non-splittable:** a major concern in MapReduce, since large files are split up into chunks across the cluster.

# Configuring Data Compression

```
Configuration conf = job.getConfiguration();
conf.setBoolean(MRJobConfig.MAP_OUTPUT_COMPRESS,
true);
conf.setClass(MRJobConfig.MAP_OUTPUT_COMPRESS_CODEC,
              SnappyCodec.class,
              CompressionCodec.class);

conf.setBoolean(FileOutputFormat.COMPRESS, true);
conf.setClass(FileOutputFormat.COMPRESS_CODEC,
              SnappyCodec.class,
              CompressionCodec.class);
```