


Topics Covered

- About Pig
- Pig Latin
- The Grunt Shell
- *Demo: Understanding Pig*
- Pig Latin Relation Names and Field Names
- Pig Data Types
- Defining a Schema
- *Lab: Getting Started with Pig*
- The GROUP Operator
- *Lab: Exploring Data with Pig*

About Pig

- It is an engine for executing programs on top of Hadoop
- It provides a language, Pig Latin, to specify these programs



Pig Latin

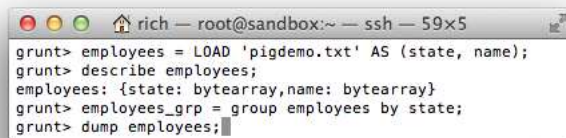
- High-level data-flow scripting language
- Pig executes in a unique fashion:
 - During execution, each statement is processed by the Pig interpreter
 - If a statement is valid, it gets added to a **logical plan** built by the interpreter
 - The steps in the logical plan do not actually execute until a DUMP or STORE command is used

The Grunt Shell

- Is an interactive shell for entering Pig Latin statements
- Is started by running the **pig** executable



Grunt shell

A terminal window titled 'rich — root@sandbox:~ — ssh — 59x5' showing the following commands and output:

```
grunt> employees = LOAD 'pigdemo.txt' AS (state, name);
grunt> describe employees;
employees: {state: bytearray,name: bytearray}
grunt> employees_grp = group employees by state;
grunt> dump employees;
```

Demo: Understanding Pig

Pig Latin Relation Names

- A **relation** is the result of a processing step
- The name given to a relation is called an **alias**
- For example, **stocks** is an alias:

```
stocks = LOAD 'mydata.txt'  
         USING TextLoader();
```

Pig Latin Field Names

- Relations can define and use field names, which are associated with an alias
- For example:

```
salaries = LOAD 'salary.data'  
           USING PigStorage(',')  
           AS (gender, age, income, zip);  
highsalaries = FILTER salaries BY income >  
1000000;
```


Pig Data Types

- `int`
- `long`
- `float`
- `double`
- `chararray`
- `bytearray`
- `boolean`
- `datetime`
- `bigdecimal`
- `biginteger`

Pig Complex Types

- **Tuple:** ordered set of values
(OH,Mark,Twain,31225)
- **Bag:** unordered collection of tuples
{
 (OH,Mark,Twain,31225),
 (UK,Charles,Dickens,42207),
 (ME,Robert,Frost,11496)
}
- **Map:** collection of key/value pairs
[state#OH,name#Mark Twain,zip#31225]

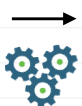
Defining a Schema

```
customers = LOAD 'customer_data' AS (  
  firstname: chararray,  
  lastname:chararray,  
  house_number:int,  
  street:chararray,  
  phone:long,  
  payment:double);
```

```
salaries = LOAD 'salaries.txt' AS  
(gender:chararray,  
  details:bag(  
    (age:int,salary:double,zip:long)  
  ));
```

Lab: Getting Started with Pig


The GROUP Operator

salaries					salariesbyage	
gender	age	salary	zip		group	salaries
F	17	41000.00	95103		17	{{(F,17,41000.0,95103), (M,17,35000.0,95103)}}
M	19	76000.00	95102		19	{{(M,19,76000.0,95102), (F,19,60000.0,95105), (M,19,14000.0,95102)}}
F	22	95000.00	95103			
F	19	60000.00	95105			
M	19	14000.00	95102			
M	17	35000.00	95103		22	{{(F,22,95000.0,95103)}}

salariesbyage = **GROUP** salaries **BY** age;

```
grunt> DESCRIBE salariesbyage;
salariesbyage: {group:int,
  salaries:{{(gender: chararray, age: int,salary: double,zip: int)}}
```


GROUP ALL

salaries					allsalaries	
gender	age	salary	zip		group	salaries
F	17	41000.00	95103		all	{(F,17,41000.0,95103), (M,19,76000.0,95102), (F,22,95000.0,95103), (F,19,60000.0,95105), (M,19,14000.0,95102), (M,17,35000.0,95103)}
M	19	76000.00	95102			
F	22	95000.00	95103			
F	19	60000.00	95105			
M	19	14000.00	95102			
M	17	35000.00	95103			

allsalaries = **GROUP** salaries **ALL**;

```
grunt> DESCRIBE allsalaries;
allsalaries: {
  group: chararray,
  salaries: {(gender: chararray,age: int,salary: double,zip: int)}
```



Relations without a Schema

salaries					salariesgroup	
\$0	\$1	\$2	\$3		group	salaries
F	17	41000.00	95103	→ 	95103	{(F,17,41000.0,95103), (F,22,95000.0,95103) (M,17,35000.0,95103)}
M	19	76000.00	95102			
F	22	95000.00	95103		95102	{(M,19,76000.0,95102), (M,19,14000.0,95102)}
F	19	60000.00	95105			
M	19	14000.00	95102			
M	17	35000.00	95103		95105	{(F,19,60000.0,95105)}

salariesgroup = **GROUP** salaries **BY** \$3;

```
grunt> DESCRIBE salariesgroup;
salariesgroup: {group:bytearray,
  salaries:{()}}
```

The FOREACH...GENERATE Operator

salaries					A	
gender	age	salary	zip		age	salary
M	66	41000.00	95103	→ 	66	41000.00
M	58	76000.00	57701		58	76000.00
F	40	95000.00	95102		40	95000.00
M	45	60000.00	95105		45	60000.00
F	28	55000.00	95103		28	55000.00

A = **FOREACH** salaries **GENERATE** age, salary;

```
grunt> DESCRIBE A;  
A: {age: int, salary: double}
```


Specifying Ranges in FOREACH

```
salaries = LOAD 'salaries.txt' USING  
PigStorage(',') AS (gender:chararray,  
age:int,salary:double,zip:int);  
C = FOREACH salaries GENERATE age..zip;  
D = FOREACH salaries GENERATE age..;  
E = FOREACH salaries GENERATE ..salary;
```

```
customer = LOAD 'data/customers';  
F = FOREACH customer GENERATE $12..$23;
```

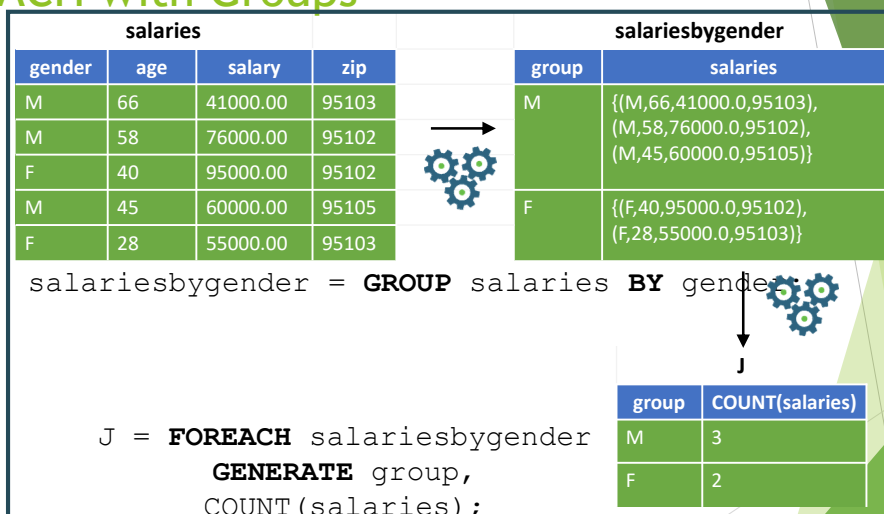
Field Names in FOREACH

```
salaries = LOAD 'salaries.txt' USING  
PigStorage(',') AS (gender:chararray,  
age:int,salary:double,zip:int);  
C = FOREACH salaries GENERATE zip, salary;  
C: {zip: int,salary: double}
```


```
D = FOREACH salaries GENERATE zip,  
    salary * 0.10;  
D: {zip: int,double}
```

```
E = FOREACH salaries GENERATE zip,  
    salary * 0.10 AS bonus;  
E: {zip: int,bonus: double}
```

FOREACH with Groups



The FILTER Operator

salaries					G			
gender	age	salary	zip		gender	age	salary	zip
F	17	41000.00	95103	→ 	M	19	76000.0	95102
M	19	76000.00	95102		F	22	95000.0	95103
F	22	95000.00	95103		F	19	60000.0	95105
F	19	60000.00	95105					
M	19	14000.00	95102					
M	17	35000.00	95103					

G = **FILTER** salaries **BY** salary >= 50000.0;

The LIMIT Operator

```
employees = LOAD 'pigdemo.txt' AS  
(state:chararray, name:chararray);  
  
emp_group = GROUP employees BY state;  
  
L = LIMIT emp_group 3;
```

Lesson Review

1. List two Pig commands that cause a logical plan to execute.
2. Which Pig command stores the output of a relation into a folder in HDFS?

```
XFR,2004-05-13,22.90,400  
XFR,2004-05-12,22.60,400000  
XFR,2004-05-11,22.80,2600  
XFR,2004-05-10,23.00,3800  
XFR,2004-05-07,23.55,2900  
XFR,2004-05-06,24.00,2200
```

```
prices = load 'prices.csv' using PigStorage(',')  
as (symbol:chararray, date:chararray, price:double, volume:int);
```

Explain what each of the following Pig commands or relations do:

3. describe prices;
4. A = group prices by symbol;
5. B = foreach prices generate symbol as x, volume as y;
6. C = foreach A generate group, SUM(prices.volume);
7. D = foreach prices generate symbol..price;

Lab: Exploring Data with Pig